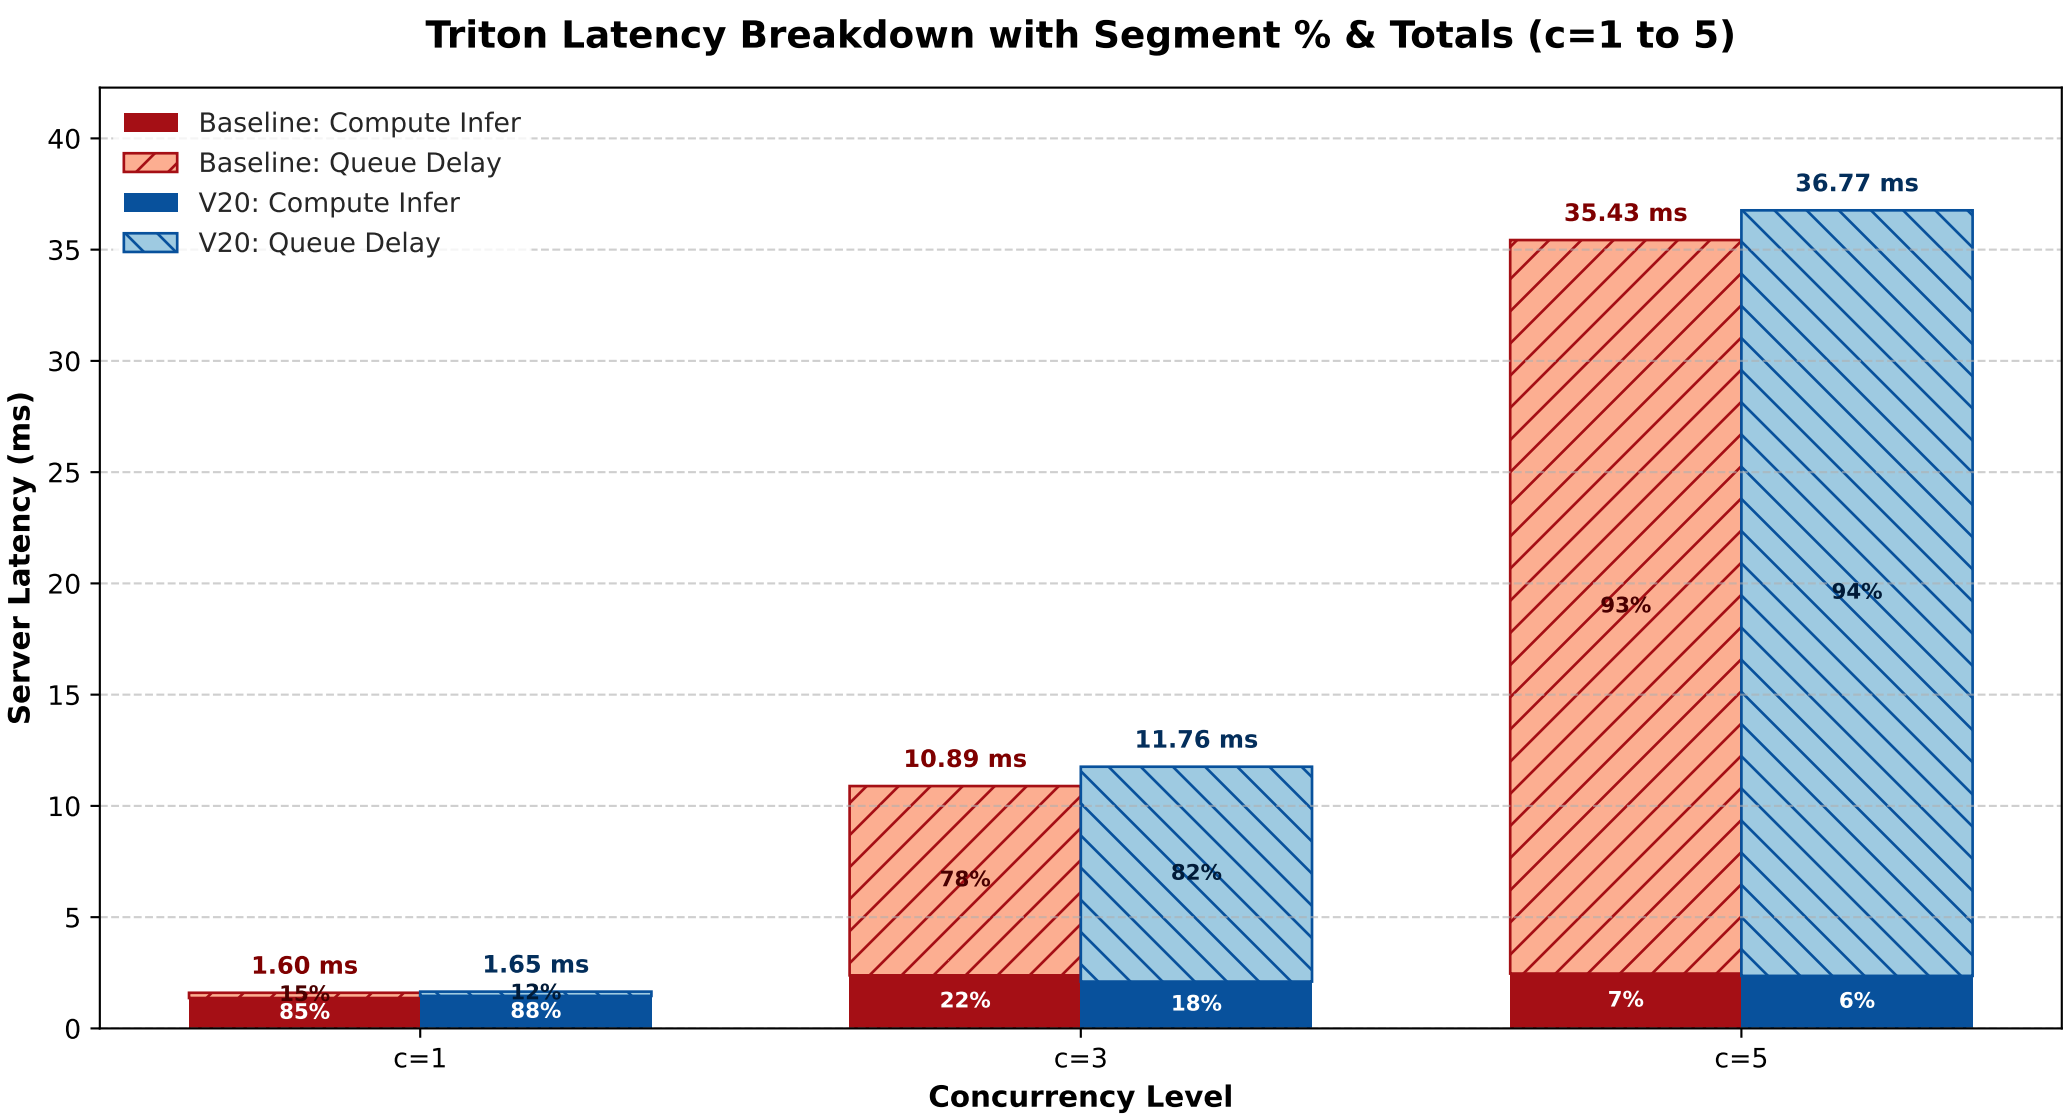
Triton Latency Breakdown with Segment % & Totals (c=1 to 5)
Baseline: Compute Infer
Baseline: Queue Delay
V20: Compute Infer
V20: Queue Delay
Server Latency (ms)
0
5
10
15
20
25
30
35
40
1.60 ms
15%
85%
1.65 ms
12%
88%
10.89 ms
78%
22%
11.76 ms
82%
18%
35.43 ms
93%
7%
36.77 ms
94%
6%
c=1
c=3
c=5
Concurrency Level